



NAME: ELEOGU CHUKWUEBUKA JOSEPH

COURSE: Electrical/electronics engineering

MATRIC NUMBER: 23120211012

P016: Voltage Stability Model Project report

GET 307: Introduction to Artificial Intelligence, Machine learning and Convergent Technologies

Department of Electrical/electronics Engineering
Pan Atlantic University



DEDICATION

First of all I would like to dedicate this project to Almighty God for the wisdom, the grace, and the perseverance granted to me throughout my studies and throughout the course of this project. I also dedicate this work to my family, especially my parents and extended family and friends for their unwavering support, financial sacrifices and continuous encouragement throughout the course of me pursuing my engineering degree and also throughout the course of this project.

ACKNOWLEDGEMENTS

I would like to, in a very special way, express my profound gratitude to my course lecturer and project supervisor, **Engr. Abdulbasit Makinde**, for his invaluable guidance, his constructive feedback and technical insights that were given to me (directly or indirectly) throughout the course of carrying out this project. His expertise and lectures in class really went a long way in me being able to say that I was able to start and complete this project.

Also, my sincere appreciation goes to the **Faculty of Electrical/Electronics Engineering** here at Pan Atlantic University for the variety of resources that i have been exposed to throughout the course of this degree (still ongoing) as well as their relentless insistence on academic excellence and hard work have driven me to see this project through to completion from start to finish.

FORMAT FOR THIS REPORT

This report was written with:

1. Font: Times New Roman
2. 12pt font size
3. A line spacing of 1.5
4. Margins (2.5cm on all sides)
5. This report will be submitted in **PDF** format.

Department of electrical/electronics

Engineering



WHAT THIS REPORT WILL CONTAIN

This project report is going to contain:

1. A title page (which I assume to be the same one as in the proposal)
2. An abstract: about a 200 to 250 word long summary that will contain what was done, how it was done, and some of the observations I've made throughout the process of conducting this project.
3. An introduction part will discuss the background, the literature review, methodological approaches that I took to conduct this research project (including the methods of how I did it), results of my efforts, including justification for the engineering choices I had to make as well as some graphs.
4. A Discussion part is going to discuss the results of this work, including model charts, performance metrics, and other findings I made during the execution of this project. It will include my analysis of the obtained results in terms of what it means from an engineering point of view.
5. A conclusion, which is going to summarize the achievements I've made during the process of this project, outline the limitations and weaknesses of the chosen approach, and provide some suggestions for further work in this area.
6. A bibliography list, which is going to include the citations to the articles and resources I had to use in order to complete this project report. They will include web-sites with ideas on how I could solve the many problems I've met when conducting this project.
7. And an appendix section will hold the actual source code of the project, some raw data samples, and additional charts/figures, if necessary.



OBJECTIVES OF THIS PROJECT

Some of the objectives that I wanted to achieve while modeling included:

1. Training a Decision Tree classifier model on Machine Learning and using it to predict grid stability and analyzing its prediction accuracy through accuracy, ROC (Receiver Operating Characteristic Curve), and AUC (Area Under Curve).

What is also essential to note is that the AUC score is ideal when it achieves 1.0, and any value below 0.5 means that it is basically either a guessing algorithm or not very good at all.

2. Analyzing and interpreting the key critical electrical factors that impact stability predictions.

3. Comparing the accuracy score of both models and analyzing and justifying their accuracy scores.

4. Developing a practical operational threshold that can help in predicting a shutdown due to instability for NERC

ABSTRACT OF THE REPORT

This project investigates the application of machine learning (ML) in order to predict voltage stability in various electrical power systems, which happens to be a critical binary classification problem as pertains to national grid security and stability.

Having utilized the UCI Electrical Grid Stability Dataset, this study evaluates the efficacy of a Logistic regression model and how it squares up against a Decision Tree classifier.

It's also important to know the differences between the two (Decision Tree Classifier and Logistical regression). One thing to know is that both of them are the two most fundamental classification algorithms in machine learning.

In machine learning, classification means trying to teach the computer how to sort data into various categories. So it happens that this model I attempted to build (and quite successfully so) is aimed at solving a classic classification problem.

When trying to build or program a machine learning model to solve classification problems, there are two main types of models that an engineer could decide to use, both solve the problem but they do so in profoundly different ways.



The dataset which comprises of about 10,000 operational instances was preprocessed to map stability states into certified binary outcomes and scaled in order to optimize gradient descent and what have you.

During the course of this project, it was discovered that the Logistical Regression Model scored an accuracy of 81.85% as opposed to the Decision Tree Classifier which scored an accuracy of 85.85%.

The Logistic regression model scored a balanced accuracy of 79.29% whilst the Decision tree classifier scored a blanched accuracy of 85.85%.

The logistic regression model scored an F1-Score of 0.7361 whilst the Decision Tree Classifier had an F1-Score of 0.8117.

The logistic regression model scored a CV Accuracy mean of 0.8154 as compared to the decision tree classifier's 0.8455

It is also important to note that the feature importance analysis revealed that producer adaptation time (τ_1) and consumer reaction times (τ_1 , τ_2 , τ_3 , τ_4) happen to be the most significant predictors that both machine learning models used for predicting grid stability.

```
— DECISION TREE FEATURE IMPORTANCE —  
Feature Importance  
3 tau4 0.153021  
2 tau3 0.146177  
1 tau2 0.145699  
0 tau1 0.128405  
8 g1 0.115032
```

In order to provide actionable operational intelligence, there was an operational threshold analysis that was conducted.

I **highly recommend** that the Nigerian Electricity Regulatory Commission (NERC) trigger emergency actions whenever the model's confidence of stability in the grid falls below 75% because this will prioritize grid safety and minimize the risk of any catastrophic failure.



INTRODUCTION TO THE REPORT

The voltage stability of a power system is the ability of the power system under normal circumstances to sustain a constant voltage profile on all the buses and transmission points on the power system over an extended operating time after being exposed to disturbances that may alter its operations.

There are many variables and different factors that can affect the performance of power grids in their daily operation. The variables include:

- Fluctuation of the load level in a power grid,
- Decentralized generation of energy output and
- Changes in reactive power margin.

Another important thing that is key for voltage stability as in the case of this project in the implementation of machine learning predictions since manual prediction approaches such as (load flow analysis) are computationally heavy and cannot be applied to emergency responses since they are very slow.

Machine learning models (like the one I successfully created) can be applied to predict grid instability within a few microseconds after analyzing complex relationships within historical and simulated data.

There are essentially two ways to go when implementing an algorithm using machine learning for tackling this challenge which may be:

1. Logistic Regression {1}:

{Logistic regression is a type of supervised machine learning {2} technique used to perform binary classifications {3}. It can, for instance, predict the likelihood of an email being classified as spam or predict if certain medical conditions exist within a person depending on the patient's test results. The logistic regression model applies the logistic (sigmoid) function to compute the probability of an event as a value between 0 and 1}

2. Decision Tree Classifier:

{Decision trees belong to the class of supervised learning {5} models that are applied in statistics {6}, data mining {7} and machine learning {8}. Using this technique, decisions are made regarding the conclusion drawn about by observations made by the model.

Department of electrical/electronics

Tree models where the target variable can take a discrete set of values are called *classification* {11} *trees*; in these tree structures, *leaves* {12} represent class labels and branches represent *conjunctions* {13} of features that lead to those class labels.


Decision trees where the target variable can take continuous values (typically *real numbers*) {14} are called *regression* {15} *trees*. More generally, the concept of regression tree can be extended to any kind of object equipped with pairwise dissimilarities such as categorical sequences]

For this particular project, I was instructed to use a logistical regression model and to ascertain its accuracy, but to also use a decision tree classifier as well, so as to be able to compare the accuracy of both types of ML classification models and to see why one is preferred over the other for this specific problem.

LITERATURE REVIEWS

The application of machine learning algorithms to power system stability issues within national grids has actually been gaining significant traction amongst countries due to the increasing complexity of smart grids.

1. <https://link.springer.com/article/10.1007/s42979-021-00463-5>:



2. <https://www.mdpi.com/1996-1073/16/22/7562>:



Open Access Article

Decentralized Smart Grid Stability Modeling with Machine Learning

by **Borna Franović**¹ , **Sandi Baressi Šegota**^{2,*} ,
Nikola Anđelić²  and **Zlatan Čar**² 

¹ HEP Group, Distribution System Operator Ltd., Viktora Cara Emina 2, 51000 Rijeka, Croatia

² Department of Automation and Electronics, Faculty of Engineering, University of Rijeka, Vukovarska 58, 51000 Rijeka, Croatia

* Author to whom correspondence should be addressed.

Energies **2023**, *16*(22), 7562; <https://doi.org/10.3390/en16227562>

Submission received: 4 October 2023 / Revised: 10 November 2023 /
Accepted: 12 November 2023 / Published: 14 November 2023

(This article belongs to the Special Issue **Sustainable/Renewable Energy Systems Analysis and Optimization**)

3. <https://www.academia.edu/71054248/>

Model Driven Approach to Smart Grid Stability Prediction in Neo4j

All Topics › Computer Science › Software Engineering

Model-Driven Approach to Smart Grid Stability Prediction in Neo4j

 **Nenad Petrović**  **Issam Al Azzoni**

2021, eNergetics 2021

 ...  7 pages

Download Free PDF

 See full PDF

 Download PDF

 Cite

 Save to Library  Share

4. **Arzamasov, V., et al. (2018). "Towards Concise Models of Grid Stability." *IEEE International Conference on Communications, Control, and Computing Technologies for Smart Grids (SmartGridComm)***. This was one of the foundational papers that introduced the simulated dataset (**UCI Electrical Grid Data**) used in this project to build this model. It was able to demonstrate that decentralized grid systems can be accurately modeled using machine learning classifiers rather than purely differential equations
5. **Kundur, P., et al. (2004). "Definition and Classification of Power System Stability." *IEEE Transactions on Power Systems***. This standard text

Department of electrical/electronics

establishes the physical definitions of voltage stability, providing the theoretical engineering framework that necessitates real-time monitoring of reactive power and voltage profiles.

6. **Amroune, M., et al. (2020). "Voltage Stability Assessment Using Artificial Neural Networks."** *Journal of Electrical Engineering & Technology*. The authors demonstrated that while Neural Networks offer high accuracy for stability prediction, their "black-box" nature makes it difficult for grid operators to extract actionable rules, highlighting the need for interpretable models like Logistic Regression.
7. **Mishra, S., & Mohanty, S. (2019). "Decision Tree-Based Security Assessment in Power Systems."** *Electric Power Systems Research*. This study explored tree-based classifiers for grid security, noting their lack of need for feature scaling, though pointing out they can be prone to overfitting in noisy grid environments.
8. **Ogunboyo, P. T., et al. (2021). "A Review of Grid Collapse Incidents in Nigeria."** *International Journal of Electrical Power & Energy Systems*. This contextual paper highlights the frequent voltage collapses in the Nigerian grid, establishing a critical gap for localized, fast-acting predictive algorithms that can integrate with NERC's operational thresholds to prevent systemic failure.

METHODOLOGY FOR BUILDING THIS MODEL FOR BUILDING THIS MODEL

Datasets: During the course of carrying out this project, the UCI Electrical Grid Stability dataset was utilized which contains about 10,000 samples. The dataset that was selected represents a 4 node electrical star system.

Out of the 10000 samples worth of data, **8000 of them samples were used to train the models**, whilst **2000 of them samples were used for testing**.

A 4 node star electrical system is one that consists of 1 producer or generator and then 3 consumers.



The data set that was selected also features reaction times (τ), nominal power balance (p), and price elasticity coefficient (g)

Tools & Platforms used to build the model:

To carry out this project, I made use of Visual studio Code Insiders, Python 3.13, Google Colab, Kaggle, Jupyter Notebooks, Wikipedia, Youtube...

Functions that the chosen tools served:

1. **Python 3.13:** This is the main language that the entire model was written in. Python is known and widely regarded for its speed, efficiency and simplicity
2. **Visual Studio Code Insiders:** This is the main IDE (**Integrated Development Environment**) that was used to write the code and run the python Machine learning environment
3. **Google colab:** This was the second IDE that I used to write my code. Google Colab is more or less like a web-based machine learning environment that contains and gives remote access to millions and millions of python libraries and dependencies without having to download them directly into your system. It saved me alot of time because i tried to download **Anaconda Navigator** which is another very powerful **IDE** just like Google colab, only that when installing, you'd definitely have to download up to 15GB worth of libraries and dependencies across multiple languages (not only python)
4. **Kaggle:** Kaggle is a platform where machine learning engineers can come together, drop datasets and also give useful tips on how to go about navigating the new and uniquely challenging world of machine learning and programming.
5. **Wikipedia:** Wikipedia is the one stop shop for articles where various concepts are explained simply and with so much clarity. Whilst writing this report, I found myself constantly being referred over to wikipedia to cite the sources for many of the definitions that are here.
6. **Youtube:** Youtube was the application where I got alot of inspiration and help from to learn some of the fundamentals of writing binary classification models, writing python scripts and programming machine learning models.



7. **Jupyter Notebooks:** Its another web based IDE used strictly for machine learning models and for advanced python scripting.

Step by step process for executing this project

1. **Download the Dataset:**

The obvious way of going about executing this project would be to first download the data set that you are going to use to train your machine learning model. In this case it's the **UCI electrical grid voltage stability data**.

2. **Load up your IDE:**

So the second obvious thing to do would be to open your IDE of choice. In my case, it's Google Colab that I first used (after I was having issues creating a python environment py -venv and installing the required dependencies on my visual studio code.) After opening your IDE, you need to download the required dependencies.

3. **Install the required dependencies for cell 1:**

So, in cell 1, we basically want to download all the dependencies that we need such as: Pandas, Seaborn, Matplotlib and pyplot.

You should have code that looks like this:

```
from matplotlib import pyplot as plt
import pandas as pd
import seaborn as sns
import os
from sklearn.model_selection import train_test_split
```

4. **Then you ensure that output directories exist and are functional:**

```
os.makedirs(CONFIG['model_dir'], exist_ok=True)
os.makedirs(CONFIG['output_dir'], exist_ok=True)

print("=" * 60)
print("DOWNLOADING DATASET")
print("=" * 60)
try:
    df = pd.read_csv(CONFIG['https://archive.ics.uci.edu/ml/machine-learning-
databases/00471/Data_for_UCI_named.csv'])
```



```
'])
    print(f"✓ Dataset downloaded successfully")
except Exception as e:
    print(f"x ERROR downloading dataset: {e}")
    raise

print("\nValidating dataset...")
if df.shape[0] == 0:
    raise ValueError("Downloaded dataset is empty!")
if 'stabf' not in df.columns:
    raise ValueError("Required column 'stabf' not found in dataset!")
print(f"✓ Dataset validation passed")

print("\n" + "=" * 60)
print("DATA PREPROCESSING & CLASS BALANCE CHECK")
print("=" * 60)
```

5. Data Preprocessing:

This involves dropping the continuous “**stab**” column in order to prevent data leakage.

The target variable “**stabf**” was converted to a binary format where 1 means that at that point in time, the voltage within the grid is stable, and where 0 means that at that point in time, the voltage is unstable. Class distribution was also checked, this revealed that there was a slight imbalance of 63.8% unstable and 36.2% stable.

6. Model Training:

A logistic regression model was then trained by using the imported data set that was then scaled.

The diagram below shows the code and the output showing the scaled version of the data:

```
print("\n" + "=" * 60)
print("DATA SCALING (for distance-based algorithms)")
print("=" * 60)
print("StandardScaler centers features to mean=0, std=1")
print("This is REQUIRED for Logistic Regression but NOT for Decision Trees\n")

scaler = StandardScaler()
X_train_scaled_arr = scaler.fit_transform(X_train)
X_test_scaled_arr = scaler.transform(X_test)
```

Department of electrical/electronics

```
# ===== ISSUE #3: Preserve feature names after scaling =====
X_train_scaled = pd.DataFrame(X_train_scaled_arr, columns=X_train.columns,
                              index=X_train.index)
X_test_scaled = pd.DataFrame(X_test_scaled_arr, columns=X_test.columns,
                              index=X_test.index)

print("✓ Scaling complete with feature names preserved")
print(f" Scaled training data shape: {X_train_scaled.shape}")
print(f" Scaled testing data shape: {X_test_scaled.shape}")
print(f" Features: {list(X_train_scaled.columns)}")
print(f" Sample scaled values (first row): {X_train_scaled.iloc[0, :3].values}")
```

The probability formula governing this particular binary classification problem is given by the diagram below.

$$P(Y = 1|X) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 X_1 + \dots + \beta_n X_n)}}$$

- A **Decision Tree** model was trained using the *unscaled* data, utilizing balanced class weights to counteract the initial dataset imbalance.

7. **Validation:** Then I proceeded to deploy a 5-Fold stratified cross validation strategy in order to ensure that the models generalized well to unseen grid scenarios.

The code below is what was used to deploy the 5-fold stratified cross validation strategy:

```
# 1. Initialize and Train the Model
model = LogisticRegression(random_state=CONFIG['random_state'], max_iter=1000)
model.fit(X_train_scaled, y_train)

print("✓ Model trained successfully")

# 2. CROSS-VALIDATION: Test generalization (K-Fold)
print(f"\n--- CROSS-VALIDATION ({CONFIG['cv_splits']}-Fold Stratified) ---")
cv_strategy = StratifiedKFold(n_splits=CONFIG['cv_splits'], shuffle=True,
                               random_state=CONFIG['random_state'])
cv_scores = cross_val_score(model, X_train_scaled, y_train, cv=cv_strategy,
                             scoring='accuracy')
cv_f1 = cross_val_score(model, X_train_scaled, y_train, cv=cv_strategy,
                         scoring='f1')
print(f"Cross-Val Accuracy: {cv_scores.mean():.4f} ± {cv_scores.std():.4f}")
```

```
print(f"Cross-Val F1-Score: {cv_f1.mean():.4f} ± {cv_f1.std():.4f}")
print(f"Individual fold accuracies: {[f'{s:.4f}' for s in cv_scores]}")

# 3. Make Predictions on the unseen 20% of data
y_pred = model.predict(X_test_scaled)
y_prob = model.predict_proba(X_test_scaled)[:, 1] # Probability of being stable
```

- 8. Evaluation:** The two main models (**The Decision Tree classifier** and the **Logistical regression model**) were evaluated and tested for their accuracy, balance accuracy, F1-Score, confusion matrices, and ROC/AUC curves (**all of which will be seen in the appendix portion of this report**)

```
# 1. Initialize the Decision Tree with CONFIG parameters
dt_model = DecisionTreeClassifier(
    max_depth=CONFIG['dt_max_depth'],
    random_state=CONFIG['random_state'],
    min_samples_split=CONFIG['dt_min_samples_split'],
    min_samples_leaf=CONFIG['dt_min_samples_leaf'],
    class_weight='balanced'
)

# 2. Train on UNSCALED data
dt_model.fit(X_train, y_train)
print("✓ Decision Tree trained on UNSCALED data")

# 3. CROSS-VALIDATION
print(f"\n--- CROSS-VALIDATION ({CONFIG['cv_splits']}-Fold Stratified) ---")
dt_cv_scores = cross_val_score(dt_model, X_train, y_train, cv=cv_strategy,
                                scoring='accuracy')
dt_cv_f1 = cross_val_score(dt_model, X_train, y_train, cv=cv_strategy,
                            scoring='f1')
print(f"Cross-Val Accuracy: {dt_cv_scores.mean():.4f} ± {dt_cv_scores.std():.4f}")
print(f"Cross-Val F1-Score: {dt_cv_f1.mean():.4f} ± {dt_cv_f1.std():.4f}")

# 4. Predictions
dt_pred = dt_model.predict(X_test)
dt_prob = dt_model.predict_proba(X_test)[:, 1]

# 5. Metrics
dt_accuracy = accuracy_score(y_test, dt_pred)
dt_balanced_acc = balanced_accuracy_score(y_test, dt_pred)
dt_f1 = f1_score(y_test, dt_pred)

# 6. Comparison Table
print("\n" + "=" * 60)
print("MODEL COMPARISON: LOGISTIC REGRESSION vs DECISION TREE")
print("=" * 60)
```

```
comparison = pd.DataFrame({
    'Metric': ['Accuracy', 'Balanced Accuracy', 'F1-Score', 'CV Accuracy Mean'],
    'Logistic Regression': [
        f"{accuracy * 100:.2f}%",
        f"{balanced_acc * 100:.2f}%",
        f"{f1:.4f}",
        f"{cv_scores.mean():.4f}"
    ],
    'Decision Tree': [
        f"{dt_accuracy * 100:.2f}%",
        f"{dt_balanced_acc * 100:.2f}%",
        f"{dt_f1:.4f}",
        f"{dt_cv_scores.mean():.4f}"
    ]
})
print("\n" + comparison.to_string(index=False))

# Best model determination
if accuracy > dt_accuracy:
    best_model = "Logistic Regression"
    best_accuracy = accuracy
else:
    best_model = "Decision Tree"
    best_accuracy = dt_accuracy

print(f"\n✓ BEST MODEL: {best_model} ({best_accuracy * 100:.2f}% accuracy)")

# Plots
fig, axes = plt.subplots(2, 2, figsize=(16, 12))

cm_lr = confusion_matrix(y_test, y_pred)
sns.heatmap(cm_lr, annot=True, fmt='d', cmap='Blues', ax=axes[0, 0])
axes[0, 0].set_title(f'Logistic Regression (Acc: {accuracy*100:.2f}%)')

cm_dt = confusion_matrix(y_test, dt_pred)
sns.heatmap(cm_dt, annot=True, fmt='d', cmap='Oranges', ax=axes[0, 1])
axes[0, 1].set_title(f'Decision Tree (Acc: {dt_accuracy*100:.2f}%)')

fpr_lr, tpr_lr, _ = roc_curve(y_test, y_prob)
axes[1, 0].plot(fpr_lr, tpr_lr, label=f'LR (AUC = {auc(fpr_lr, tpr_lr):.3f})')
axes[1, 0].set_title('ROC - Logistic Regression')

fpr_dt, tpr_dt, _ = roc_curve(y_test, dt_prob)
axes[1, 1].plot(fpr_dt, tpr_dt, color='green', label=f'DT (AUC = {auc(fpr_dt, tpr_dt):.3f})')
axes[1, 1].set_title('ROC - Decision Tree')

plt.tight_layout()
plt.show()
```



```
# Feature Importance
dt_feature_importance = pd.DataFrame({
    'Feature': X.columns,
    'Importance': dt_model.feature_importances_
}).sort_values('Importance', ascending=False)

print("\n--- DECISION TREE FEATURE IMPORTANCE ---")
print(dt_feature_importance.head(5))

dt_importance_path = os.path.join(CONFIG['output_dir'],
    'feature_importance_decision_tree.csv')
dt_feature_importance.to_csv(dt_importance_path, index=False)
```

```
best_accuracy = dt_accuracy

print(f"\n✓ BEST MODEL: {best_model} ({best_accuracy * 100:.2f}% accuracy)")

# Plots
fig, axes = plt.subplots(2, 2, figsize=(16, 12))

cm_lr = confusion_matrix(y_test, y_pred)
sns.heatmap(cm_lr, annot=True, fmt='d', cmap='Blues', ax=axes[0, 0])
axes[0, 0].set_title(f'Logistic Regression (Acc: {accuracy*100:.2f}%)')

cm_dt = confusion_matrix(y_test, dt_pred)
sns.heatmap(cm_dt, annot=True, fmt='d', cmap='Oranges', ax=axes[0, 1])
axes[0, 1].set_title(f'Decision Tree (Acc: {dt_accuracy*100:.2f}%)')

fpr_lr, tpr_lr, _ = roc_curve(y_test, y_prob)
axes[1, 0].plot(fpr_lr, tpr_lr, label=f'LR (AUC = {auc(fpr_lr, tpr_lr):.3f})')
axes[1, 0].set_title('ROC - Logistic Regression')

fpr_dt, tpr_dt, _ = roc_curve(y_test, dt_prob)
axes[1, 1].plot(fpr_dt, tpr_dt, color='green', label=f'DT (AUC = {auc(fpr_dt,
tpr_dt):.3f})')
axes[1, 1].set_title('ROC - Decision Tree')

plt.tight_layout()
plt.show()

# Feature Importance
dt_feature_importance = pd.DataFrame({
    'Feature': X.columns,
    'Importance': dt_model.feature_importances_
}).sort_values('Importance', ascending=False)

print("\n--- DECISION TREE FEATURE IMPORTANCE ---")
print(dt_feature_importance.head(5))
```

Department of electrical/electronics

Engineering


```
dt_importance_path = os.path.join(CONFIG['output_dir'],  
'feature_importance_decision_tree.csv')  
dt_feature_importance.to_csv(dt_importance_path, index=False)
```

RESULTS & DISCUSSIONS

1. Logistic Regression Performance:

The logistic regression model achieved an overall accuracy of **81.85%**. Also in the output section that compared the accuracy, F1 score for the decision tree classifier and the logistic regression models the precision for predicting unstable states was high (0.85) which indicated that the Logistical regression model that i developed demonstrated a strong **True Positive Rate** against **False Positive Rate**, yielding a highly capable AUC (Area Under Curve) score.

2. Feature Importance Analysis:

Whilst analysing the absolute coefficients of the **logistic regression model**, I discovered that there were two main variables that caused about 70-80% of grid instability. They happened to be:

- I. □2 (consumer1 reaction time): This had one of the highest negative coefficients of the lot. According to the output cell 4, **□2 has a negative coefficient of -0.8867** and an **impact level of 0.8867**. It's important to note that the values range between 1 and -1
- II. □1 (producer adaptation time): This feature had a **negative coefficient of -0.8910** and an **impact level of 0.8910**.
- III. □3 (consumer 2 reaction time): This feature had a **negative coefficient of -0.8657** and an **impact level of 0.8657**.

However for positive coefficients, we have **p4, p2 and p1**. These are the variables within the electrical grid that are **inherently stabilizing in nature**.

Also, this model recorded an **Intercept bias of -1.0558**.

Comparison between the Logistic Regression model and the decision tree classifier

Department of electrical/electronics



On the output for cell 5, the machine learning algorithm compared and contrasted the two main models that I had run in different cells. Both the **decision tree classification model** and the **Logistic regression model** were tried and tested, and after that i compared the **CV average mean** and decided to go with the model with the significantly higher mean accuracy which was the **decision tree model**.

The screenshot taken from the IDE below helps confirm the output which was received:

```
=====
DECISION TREE MODEL - COMPARISON & EVALUATION
=====
🔑 KEY DIFFERENCE: Decision Trees work on UNSCALED data (no scaling needed)

✓ Decision Tree trained on UNSCALED data

--- CROSS-VALIDATION (5-Fold Stratified) ---
Cross-Val Accuracy: 0.8455 ± 0.0116
Cross-Val F1-Score: 0.7960 ± 0.0146

=====
MODEL COMPARISON: LOGISTIC REGRESSION vs DECISION TREE
=====

Metric Logistic Regression Decision Tree
Accuracy          81.85%          85.85%
Balanced Accuracy  79.23%          85.50%
F1-Score          0.7364          0.8117
CV Accuracy Mean  0.8154          0.8455

✓ BEST MODEL: Decision Tree (85.85% accuracy)
```

From the screenshot above, it is clear that the **Decision Tree Model** outperformed the **Logistic regression model**. And I believe the major reason for this happening is that the **relationship between the mechanical/electrical grid parameters and ultimate voltage stability hardly has any underlying linear correlations**.

Power systems are highly non linear in nature and the voltage stability will depend on many complex interactions between parameters like reactive power limits, rotor angles, load variance and so on.

What I mean is that, there is hardly any linear relationship between values and the states of the voltage stability within a system. So maybe after the tau values exceed a certain arithmetic figure, we have a situation where the voltage level within the electrical grid system might become unstable, and whilst it's within a certain range, it remains stable.

So there seems to be a **nonlinear** relationship between the electrical grid values and the voltage stability across the system (grid)



PRACTICAL RECOMMENDATIONS GOING INTO THE FUTURE

So even though it's not common knowledge, it's good to know that most of the Voltage stability prediction models assume a 50% threshold for binary classification, but this brings about an issue.

In power systems, a **false negative (like missing a grid collapsing voltage spike)** is actually going to be catastrophic (financially and structurally in terms of damage). But however, a **false positive (unnecessary load/voltage shedding)** is merely an inconvenience which might even be a good thing (arguably). So based on the operational threshold analysis that was carried out in the code below:

```
1. Get the raw probabilities of being STABLE (Class 1)
probabilities = model.predict_proba(X_test_scaled)[: , 1]

# 2. Test multiple thresholds to find optimal trade-off
thresholds_to_test = [0.50, 0.60, 0.70, 0.75, 0.80, 0.90]

print("\nThreshold Analysis - Testing Different Safety Levels:\n")
print(f"{'Threshold':<12} {'Sensitivity':<12} {'Specificity':<12} {'Precision':<12} {'F1-Score':<12}")
print("-" * 60)

threshold_results = []

for threshold in thresholds_to_test:
    custom_pred = (probabilities >= threshold).astype(int)

    # Calculate metrics
    tn, fp, fn, tp = confusion_matrix(y_test, custom_pred).ravel()

    sensitivity = tp / (tp + fn) if (tp + fn) > 0 else 0 # True Positive Rate
    specificity = tn / (tn + fp) if (tn + fp) > 0 else 0 # True Negative Rate
    precision = tp / (tp + fp) if (tp + fp) > 0 else 0
    fl_th = 2 * (precision * sensitivity) / (precision + sensitivity) if (precision +
sensitivity) > 0 else 0

    threshold_results.append({
        'Threshold': threshold,
        'Sensitivity': sensitivity,
        'Specificity': specificity,
        'Precision': precision,
        'F1': fl_th,
        'Custom_Pred': custom_pred
    })

    print(f"{'threshold':<12.2f} {'sensitivity':<12.3f} {'specificity':<12.3f}
{'precision':<12.3f} {'fl_th':<12.3f}")
```

Department of electrical/electronics



```
# 3. Recommend optimal threshold
optimal_idx = np.argmax([r['F1'] for r in threshold_results])
optimal_threshold = threshold_results[optimal_idx]['Threshold']
print(f"\n✓ RECOMMENDED THRESHOLD: {optimal_threshold:.2f} (F1-Score: {threshold_results[optimal_idx]['F1']:.4f})")

# ===== ISSUE #4: Use configurable NERC threshold from CONFIG =====
nerc_threshold = CONFIG['nerc_threshold']
custom_predictions = (probabilities >= nerc_threshold).astype(int)

print(f"\n" + "=" * 60)
print(f"NERC EARLY WARNING SYSTEM (Threshold = {nerc_threshold*100:.0f}%)")
print("=" * 60)
print(f"Interpretation: Grid must have ≥75% confidence to be declared SAFE")
print(f"Otherwise: EMERGENCY ACTION TRIGGERED\n")

nerc_cm = confusion_matrix(y_test, custom_predictions)
tn, fp, fn, tp = nerc_cm.ravel()

print("Classification Report:")
print(classification_report(y_test, custom_predictions,
                           target_names=['Emergency Action', 'Grid Safe']))

print(f"\nOperational Breakdown:")
print(f" True Negatives (Correct Emergency): {tn:5d} (avoiding false alarms)")
print(f" False Positives (False Alarms): {fp:5d} (unnecessary activations)")
print(f" False Negatives (MISSED RISKS): {fn:5d} ⚠ CRITICAL")
print(f" True Positives (Correct Safe): {tp:5d}")

# 5. Visualize threshold impact
fig, axes = plt.subplots(1, 2, figsize=(15, 5))

# Plot 1: Metrics vs Threshold
thresholds_plot = [r['Threshold'] for r in threshold_results]
sensitivity_plot = [r['Sensitivity'] for r in threshold_results]
specificity_plot = [r['Specificity'] for r in threshold_results]
f1_plot = [r['F1'] for r in threshold_results]

axes[0].plot(thresholds_plot, sensitivity_plot, marker='o', label='Sensitivity (TP Rate)', linewidth=2)
axes[0].plot(thresholds_plot, specificity_plot, marker='s', label='Specificity (TN Rate)', linewidth=2)
axes[0].plot(thresholds_plot, f1_plot, marker='^', label='F1-Score', linewidth=2)
axes[0].axvline(x=nerc_threshold, color='red', linestyle='--', linewidth=2, label=f'NERC Threshold ({nerc_threshold})')
axes[0].set_xlabel('Probability Threshold')
axes[0].set_ylabel('Score')
axes[0].set_title('Threshold Optimization Analysis')
axes[0].legend()
```

Department of electrical/electronics

Engineering



```
axes[0].grid(True, alpha=0.3)

# Plot 2: NERC Confusion Matrix
sns.heatmap(nerc_cm, annot=True, fmt='d', cmap='RdYlGn',
            xticklabels=['Emergency', 'Safe'],
            yticklabels=['Actually Unstable', 'Actually Stable'],
            ax=axes[1], cbar_kws={'label': 'Count'})
axes[1].set_title(f'NERC Threshold System (Threshold={nerc_threshold})')
axes[1].set_ylabel('True Label')
axes[1].set_xlabel('System Decision')

plt.tight_layout()
plt.show()
```

The model was able to compute the theoretical probability of being **Stable**, and then carrying out **fresh-hold testing** in order to find the optimal tradeoff.

After that was done, the metrics were calculated using **confusion matrices** (all of these are in the code uploaded above).

After the model would have carried out **fresh-hold testing** and **confusion matrix calculations** then it will generally come up with a recommended optimal threshold after you run this code:

```
3. Recommend optimal threshold
optimal_idx = np.argmax([r['F1'] for r in threshold_results])
optimal_threshold = threshold_results[optimal_idx]['Threshold']
print(f"\n✓ RECOMMENDED THRESHOLD: {optimal_threshold:.2f} (F1-Score:
{threshold_results[optimal_idx]['F1']:.4f})")
```

Now this is where we use a configurable NERC threshold from Config:

```
nerc_threshold = CONFIG['nerc_threshold']
custom_predictions = (probabilities >= nerc_threshold).astype(int)

print(f"\n" + "=" * 60)
print(f"NERC EARLY WARNING SYSTEM (Threshold = {nerc_threshold*100:.0f}%)")
print("=" * 60)
print("Interpretation: Grid must have ≥75% confidence to be declared SAFE")
print("Otherwise: EMERGENCY ACTION TRIGGERED\n")

nerc_cm = confusion_matrix(y_test, custom_predictions)
tn, fp, fn, tp = nerc_cm.ravel()

print("Classification Report:")
print(classification_report(y_test, custom_predictions,
                           target_names=['Emergency Action', 'Grid Safe']))

print(f"\nOperational Breakdown:")
print(f" True Negatives (Correct Emergency): {tn:5d} (avoiding false alarms)")
```

Department of electrical/electronics

Engineering



```
print(f" False Positives (False Alarms):      {fp:5d} (unnecessary activations)")
print(f" False Negatives (MISSED RISKS):      {fn:5d} ⚠ CRITICAL")
print(f" True Positives (Correct Safe):        {tp:5d}")
```

(All this code that i am uploading already happens to be in the larger upload that i made at the start of this paragraph).

Anyways, all this aside, the main recommendation that im making is that the NERC (Nigerian Electricity Regulatory Commission **set its new threshold to 0.75 or 75%**)

This means that the AI needs to be **at least 75% confident** that the grid voltages are stable before commencing normal operations.

If the model's stability prediction confidence level falls below 75%, there should be some sort of warning or **an Emergency Action** be triggered.

CONCLUSION

In this regard, this particular project managed to create a successful predictive model based on AI for voltage stability in power systems. As it can be seen from the results of using the logistic regression algorithm, its performance in creating a predictive model during this particular project is extremely reliable, as it managed to achieve a remarkable score of 81.85%, which is much worse than that of decision tree classifier with an accuracy of 85.85%.

According to the results of my analysis done in various cells of this particular IDE workspace, reaction times (τ) inside the grid turned out to be the key reasons for such kind of instability in the electrical grid, as I mentioned before. In order to improve the confidence level of this particular model, it is suggested to increase the decision threshold up to 75%.

Limitations to this model:

The major limitation that this model has is that **it is designed primarily on a simulated 4-node grid that contains 1 generator and 3 primary receivers or consumers**. And this model is severely ill-equipped for the actual topological configuration of Nigeria's electrical grid system.



So the obvious recommendation for anyone audacious enough to even think about it would be to design a voltage stability prediction model for a grid system designed for millions and millions of nodes all across the federal Republic of Nigeria. data that had been scaled using part of the code shown above. This is the phase where the algorithm proceeds to use the 2000 samples of labelled data and attempts to then correctly label the remaining 8000 samples of unsampled data with a high enough accuracy level.

REFERENCES USED WHILST WRITING THIS REPORT

1. https://en.wikipedia.org/wiki/Logistic_regression *Logistic regression*. (n.d.). Wikipedia. Retrieved May 27, 2026.
2. https://en.wikipedia.org/wiki/Supervised_learning *supervised learning* (in the definition of Logistical regression). Wikipedia. Retrieved May 27, 2026.
3. https://en.wikipedia.org/wiki/Binary_classification *binary classification* (also in the definition of logistic regression). Wikipedia. Retrieved May 27, 2026.
4. https://en.wikipedia.org/wiki/Supervised_learning *Supervised learning* Wikipedia. Retrieved May 27, 2026.
5. <https://archive.ics.uci.edu/dataset/471/electrical+grid+stability+simulated+data> *UCI Electrical Grid Simulated multivariate Data* (the main dataset used to prime/build the machine learning model). Retrieved May 27, 2026.
6. <https://en.wikipedia.org/wiki/Statistics> *statistics* Wikipedia. Retrieved May 27, 2026.
7. https://en.wikipedia.org/wiki/Data_mining *data mining*. Wikipedia. Retrieved May 27, 2026.
8. https://en.wikipedia.org/wiki/Machine_learning *machine learning*. Wikipedia. Retrieved May 27, 2026.
9. https://en.wikipedia.org/wiki/Decision_tree *decision trees*. Wikipedia. Retrieved May 27, 2026.
10. https://en.wikipedia.org/wiki/Predictive_modelling *Predictive modelling*. Wikipedia. Retrieved May 27, 2026.



11. https://en.wikipedia.org/wiki/Statistical_classification *statistical classification*. Wikipedia. Retrieved May 27, 2026.
12. [https://en.wikipedia.org/wiki/Tree_\(abstract_data_type\)#Terminology](https://en.wikipedia.org/wiki/Tree_(abstract_data_type)#Terminology) *trees*. Wikipedia. Retrieved May 27, 2026.
13. https://en.wikipedia.org/wiki/Logical_conjunction *logical conjunction*. Wikipedia. Retrieved May 27, 2026.
14. https://en.wikipedia.org/wiki/Real_number *real numbers*. Wikipedia. Retrieved May 27, 2026.
15. https://en.wikipedia.org/wiki/Regression_analysis *regression analysis*. Wikipedia. Retrieved May 27, 2026

APPENDIX SECTION

1. **APPENDIX A:** The full annotated source code:

Cell 1:

```
from matplotlib import pyplot as plt

import pandas as pd
import seaborn as sns
import os
from sklearn.model_selection import train_test_split

# Define configuration parameters
CONFIG = {
    'model_dir': './models',
    'output_dir': './outputs',
    'dataset_url':
'https://archive.ics.uci.edu/ml/machine-learning-databases/00471/Data_for_UCI_n
amed.csv',
    'test_size': 0.2,
    'random_state': 42,
    'cv_splits': 5,
    'nerc_threshold': 0.75, # NERC threshold for stability (75% confidence for
Grid Safe)
    'dt_max_depth': 10,
    'dt_min_samples_split': 2,
    'dt_min_samples_leaf': 1
}

# Ensure output directories exist
```




```
os.makedirs(CONFIG['model_dir'], exist_ok=True)
os.makedirs(CONFIG['output_dir'], exist_ok=True)

print("=" * 60)
print("DOWNLOADING DATASET")
print("=" * 60)
try:
    df = pd.read_csv(CONFIG['dataset_url'])
    print(f"✓ Dataset downloaded successfully")
except Exception as e:
    print(f"✗ ERROR downloading dataset: {e}")
    raise

print("\nValidating dataset...")
if df.shape[0] == 0:
    raise ValueError("Downloaded dataset is empty!")
if 'stabf' not in df.columns:
    raise ValueError("Required column 'stabf' not found in dataset!")
print(f"✓ Dataset validation passed")

print("\n" + "=" * 60)
print("DATA PREPROCESSING & CLASS BALANCE CHECK")
print("=" * 60)

# Drop continuous 'stab', focus on 'stabf'
df = df.drop(columns=['stab'])
df['stabf'] = df['stabf'].map({'stable': 1, 'unstable': 0})

print("\n✓ Class Distribution BEFORE splitting:")
class_counts = df['stabf'].value_counts()
class_pct = df['stabf'].value_counts(normalize=True) * 100
for label, count in class_counts.items():
    pct = class_pct[label]
    status = "Stable" if label == 1 else "Unstable"
    print(f" {status:10} (Class {label}): {count:6} samples ({pct:5.2f}%)")

X = df.drop(columns=['stabf'])
y = df['stabf']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=CONFIG['test_size'], random_state=CONFIG['random_state'],
    stratify=y
)

print(f"Training set size: {X_train.shape[0]} samples")
```

Department of electrical/electronics

Engineering



```
print(f"Testing set size: {X_test.shape[0]} samples")

print("\n" + "=" * 60)
print("FEATURE CORRELATION WITH STABILITY")
print("=" * 60)
correlation_with_stability = X.corrwith(y).sort_values(ascending=False)
print("\nCorrelation with Stability (1=Stable, 0=Unstable):")
print(correlation_with_stability)

plt.figure(figsize=(12, 8))
correlation_matrix = pd.concat([X, y], axis=1).corr()
sns.heatmap(correlation_matrix, annot=True, fmt='.2f', cmap='coolwarm',
            center=0, cbar_kws={'label': 'Correlation'})
plt.title('Feature Correlation Matrix with Stability Target')
plt.tight_layout()
plt.show()
```

Cell 2:

```
from sklearn.preprocessing import StandardScaler

print("\n" + "=" * 60)
print("DATA SCALING (for distance-based algorithms)")
print("=" * 60)
print("StandardScaler centers features to mean=0, std=1")
print("This is REQUIRED for Logistic Regression but NOT for Decision Trees\n")

scaler = StandardScaler()
X_train_scaled_arr = scaler.fit_transform(X_train)
X_test_scaled_arr = scaler.transform(X_test)

# ===== ISSUE #3: Preserve feature names after scaling =====
X_train_scaled = pd.DataFrame(X_train_scaled_arr, columns=X_train.columns,
                              index=X_train.index)
X_test_scaled = pd.DataFrame(X_test_scaled_arr, columns=X_test.columns,
                             index=X_test.index)

print("✓ Scaling complete with feature names preserved")
print(f" Scaled training data shape: {X_train_scaled.shape}")
print(f" Scaled testing data shape: {X_test_scaled.shape}")
print(f" Features: {list(X_train_scaled.columns)}")
print(f" Sample scaled values (first row): {X_train_scaled.iloc[0,
:3].values}")
```



Cell 3:

```
import seaborn as sns

from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import StratifiedKFold, cross_val_score
from sklearn.metrics import accuracy_score, balanced_accuracy_score, f1_score,
classification_report, confusion_matrix, roc_curve, auc, precision_recall_curve
import matplotlib.pyplot as plt

print("\n" + "=" * 60)
print("LOGISTIC REGRESSION - TRAINING & EVALUATION")
print("=" * 60)

# 1. Initialize and Train the Model
model = LogisticRegression(random_state=CONFIG['random_state'], max_iter=1000)
model.fit(X_train_scaled, y_train)

print("✓ Model trained successfully")

# 2. CROSS-VALIDATION: Test generalization (K-Fold)
print(f"\n--- CROSS-VALIDATION ({CONFIG['cv_splits']}-Fold Stratified) ---")
cv_strategy = StratifiedKFold(n_splits=CONFIG['cv_splits'], shuffle=True,
                               random_state=CONFIG['random_state'])
cv_scores = cross_val_score(model, X_train_scaled, y_train, cv=cv_strategy,
                             scoring='accuracy')
cv_f1 = cross_val_score(model, X_train_scaled, y_train, cv=cv_strategy,
                         scoring='f1')
print(f"Cross-Val Accuracy: {cv_scores.mean():.4f} \u00b1 {cv_scores.std():.4f}")
print(f"Cross-Val F1-Score: {cv_f1.mean():.4f} \u00b1 {cv_f1.std():.4f}")
print(f"Individual fold accuracies: {[f'{s:.4f}' for s in cv_scores]}")

# 3. Make Predictions on the unseen 20% of data
y_pred = model.predict(X_test_scaled)
y_prob = model.predict_proba(X_test_scaled)[: , 1] # Probability of being
stable

# 4. Calculate Multiple Metrics
accuracy = accuracy_score(y_test, y_pred)
balanced_acc = balanced_accuracy_score(y_test, y_pred) # Better for imbalanced
data
f1 = f1_score(y_test, y_pred)
```



```
print("\n" + "=" * 60)
print("TEST SET PERFORMANCE")
print("=" * 60)
print(f"Accuracy:          {accuracy * 100:6.2f}%")
print(f"Balanced Accuracy: {balanced_acc * 100:6.2f}% (better for
imbalanced)")
print(f"F1-Score:          {f1:.4f} (harmonic mean of precision & recall)")

print("\nDetailed Classification Report:")
print(classification_report(y_test, y_pred, target_names=['Unstable (0)',
'Stable (1)']))

# 5. CONFUSION MATRIX - Visual breakdown
cm = confusion_matrix(y_test, y_pred)
print("\nConfusion Matrix:")
print(f"                Predicted")
print(f"                Unstable  Stable")
print(f"Actual Unstable  {cm[0,0]:5d}  {cm[0,1]:5d}")
print(f"Actual Stable   {cm[1,0]:5d}  {cm[1,1]:5d}")

# Confusion Matrix Visualization
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Plot 1: Confusion Matrix Heatmap
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Unstable', 'Stable'],
            yticklabels=['Unstable', 'Stable'],
            cbar_kws={'label': 'Count'}, ax=axes[0])
axes[0].set_title('Confusion Matrix - Logistic Regression')
axes[0].set_ylabel('True Label')
axes[0].set_xlabel('Predicted Label')

# Plot 2: ROC Curve
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)

axes[1].plot(fpr, tpr, color='darkorange', lw=2.5, label=f'ROC curve (AUC =
{roc_auc:.3f})')
axes[1].plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--', label='Random
Classifier')
axes[1].fill_between(fpr, tpr, alpha=0.2, color='darkorange')
axes[1].set_xlim([0.0, 1.0])
axes[1].set_ylim([0.0, 1.05])
axes[1].set_xlabel('False Positive Rate (False Alarms)')
axes[1].set_ylabel('True Positive Rate (Correct Detections)')
```

Department of electrical/electronics



```
axes[1].set_title('ROC Curve - Grid Stability Prediction')
axes[1].legend(loc="lower right")
axes[1].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

# 6. PRECISION-RECALL CURVE (Alternative to ROC)
precision, recall, _ = precision_recall_curve(y_test, y_prob)
plt.figure(figsize=(10, 6))
plt.plot(recall, precision, marker='o', label=f'PR curve (F1={f1:.3f})')
plt.xlabel('Recall (True Positive Rate)')
plt.ylabel('Precision (True Positive / Predicted Positive)')
plt.title('Precision-Recall Curve - Grid Stability')
plt.legend()
plt.grid(True, alpha=0.3)
plt.show()
```

Cell 4

```
import numpy as np

print("\n" + "=" * 60)
print("FEATURE IMPORTANCE & GRID VULNERABILITY ANALYSIS")
print("=" * 60)

# 1. Extract the coefficients (weights) from your trained model
coefficients = model.coef_[0]
feature_names = X.columns
intercept = model.intercept_[0]

# 2. Create a detailed table
feature_importance = pd.DataFrame({
    'Feature': feature_names,
    'Coefficient': coefficients,
    'Abs_Coefficient': np.abs(coefficients)
})

# 3. Sort by absolute impact
feature_importance = feature_importance.sort_values(by='Abs_Coefficient',
ascending=False)

print("\nFeature Importance Ranking:")
print("(Positive = Stabilizing | Negative = Destabilizing)\n")
```

Department of electrical/electronics



```
for idx, row in feature_importance.iterrows():
    impact_type = "STABILIZING " if row['Coefficient'] > 0 else "DESTABILIZING"
    print(f"{row['Feature']:10} | Coef: {row['Coefficient']:8.4f} | {impact_type}
| Impact: {row['Abs_Coefficient']:.4f}")

print(f"\nModel Intercept (bias): {intercept:.4f}")

# ===== ISSUE #8: Export feature importance to CSV =====
importance_path = os.path.join(CONFIG['output_dir'],
'feature_importance_logistic.csv')
feature_importance.to_csv(importance_path, index=False)
print(f"\n✓ Feature importance saved to: {importance_path}")

# 4. Visualization
fig, axes = plt.subplots(1, 2, figsize=(16, 6))

# Plot 1: Sorted bar plot
colors = ['green' if c > 0 else 'red' for c in
feature_importance['Coefficient']]
axes[0].barh(feature_importance['Feature'], feature_importance['Coefficient'],
color=colors, alpha=0.7)
axes[0].set_xlabel('Coefficient Value')
axes[0].set_title('Feature Importance (Logistic Regression)\nGreen=Stabilizing,
Red=Destabilizing')
axes[0].axvline(x=0, color='black', linestyle='--', linewidth=0.8)
axes[0].grid(True, alpha=0.3, axis='x')

# Plot 2: Absolute importance
axes[1].barh(feature_importance['Feature'],
feature_importance['Abs_Coefficient'],
color='steelblue', alpha=0.7)
axes[1].set_xlabel('Absolute Impact')
axes[1].set_title('Feature Importance by Absolute Impact\n(Higher = More
Critical to Decision)')
axes[1].grid(True, alpha=0.3, axis='x')

plt.tight_layout()
plt.show()

# 5. Top 3 Grid Parameters to Monitor
print("\n" + "=" * 60)
print(" 🚨 TOP 3 CRITICAL PARAMETERS FOR GRID STABILITY")
print("=" * 60)
top_3 = feature_importance.head(3)
for rank, (idx, row) in enumerate(top_3.iterrows(), 1):
```

Department of electrical/electronics



```
print(f"{rank}. {row['Feature']:15} (Impact: {row['Abs_Coefficient']:.4f})")
```

Cell 5

```
import pandas as pd

import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import cross_val_score
from sklearn.metrics import accuracy_score, balanced_accuracy_score, f1_score,
classification_report, confusion_matrix, roc_curve, auc
import os

print("\n" + "=" * 60)
print("DECISION TREE MODEL - COMPARISON & EVALUATION")
print("=" * 60)
print("🔑 KEY DIFFERENCE: Decision Trees work on UNSCALED data (no scaling
needed)\n")

# 1. Initialize the Decision Tree with CONFIG parameters
dt_model = DecisionTreeClassifier(
    max_depth=CONFIG['dt_max_depth'],
    random_state=CONFIG['random_state'],
    min_samples_split=CONFIG['dt_min_samples_split'],
    min_samples_leaf=CONFIG['dt_min_samples_leaf'],
    class_weight='balanced'
)

# 2. Train on UNSCALED data
dt_model.fit(X_train, y_train)
print("✓ Decision Tree trained on UNSCALED data")

# 3. CROSS-VALIDATION
print(f"\n--- CROSS-VALIDATION ({CONFIG['cv_splits']}-Fold Stratified) ---")
dt_cv_scores = cross_val_score(dt_model, X_train, y_train, cv=cv_strategy,
scoring='accuracy')
dt_cv_f1 = cross_val_score(dt_model, X_train, y_train, cv=cv_strategy,
scoring='f1')
print(f"Cross-Val Accuracy: {dt_cv_scores.mean():.4f} ±
{dt_cv_scores.std():.4f}")
print(f"Cross-Val F1-Score: {dt_cv_f1.mean():.4f} ± {dt_cv_f1.std():.4f}")

# 4. Predictions
dt_pred = dt_model.predict(X_test)
```



```
dt_prob = dt_model.predict_proba(X_test)[: , 1]

# 5. Metrics
dt_accuracy = accuracy_score(y_test, dt_pred)
dt_balanced_acc = balanced_accuracy_score(y_test, dt_pred)
dt_f1 = f1_score(y_test, dt_pred)

# 6. Comparison Table
print("\n" + "=" * 60)
print("MODEL COMPARISON: LOGISTIC REGRESSION vs DECISION TREE")
print("=" * 60)

comparison = pd.DataFrame({
    'Metric': ['Accuracy', 'Balanced Accuracy', 'F1-Score', 'CV Accuracy Mean'],
    'Logistic Regression': [
        f"{accuracy * 100:.2f}%",
        f"{balanced_acc * 100:.2f}%",
        f"{f1:.4f}",
        f"{cv_scores.mean():.4f}"
    ],
    'Decision Tree': [
        f"{dt_accuracy * 100:.2f}%",
        f"{dt_balanced_acc * 100:.2f}%",
        f"{dt_f1:.4f}",
        f"{dt_cv_scores.mean():.4f}"
    ]
})
print("\n" + comparison.to_string(index=False))

# Best model determination
if accuracy > dt_accuracy:
    best_model = "Logistic Regression"
    best_accuracy = accuracy
else:
    best_model = "Decision Tree"
    best_accuracy = dt_accuracy

print(f"\n✓ BEST MODEL: {best_model} ({best_accuracy * 100:.2f}% accuracy)")

# Plots
fig, axes = plt.subplots(2, 2, figsize=(16, 12))

cm_lr = confusion_matrix(y_test, y_pred)
sns.heatmap(cm_lr, annot=True, fmt='d', cmap='Blues', ax=axes[0, 0])
axes[0, 0].set_title(f'Logistic Regression (Acc: {accuracy*100:.2f}%)')
```

Department of electrical/electronics



```
cm_dt = confusion_matrix(y_test, dt_pred)
sns.heatmap(cm_dt, annot=True, fmt='d', cmap='Oranges', ax=axes[0, 1])
axes[0, 1].set_title(f'Decision Tree (Acc: {dt_accuracy*100:.2f}%)')

fpr_lr, tpr_lr, _ = roc_curve(y_test, y_prob)
axes[1, 0].plot(fpr_lr, tpr_lr, label=f'LR (AUC = {auc(fpr_lr, tpr_lr):.3f})')
axes[1, 0].set_title('ROC - Logistic Regression')

fpr_dt, tpr_dt, _ = roc_curve(y_test, dt_prob)
axes[1, 1].plot(fpr_dt, tpr_dt, color='green', label=f'DT (AUC = {auc(fpr_dt,
tpr_dt):.3f})')
axes[1, 1].set_title('ROC - Decision Tree')

plt.tight_layout()
plt.show()

# Feature Importance
dt_feature_importance = pd.DataFrame({
    'Feature': X.columns,
    'Importance': dt_model.feature_importances_
}).sort_values('Importance', ascending=False)

print("\n--- DECISION TREE FEATURE IMPORTANCE ---")
print(dt_feature_importance.head(5))

dt_importance_path = os.path.join(CONFIG['output_dir'],
'feature_importance_decision_tree.csv')
dt_feature_importance.to_csv(dt_importance_path, index=False)
```

Cell 6:

```
import numpy as np

import joblib # Import joblib for saving and loading models

print("\n" + "=" * 60)
print("OPERATIONAL THRESHOLD ANALYSIS - NERC EARLY WARNING SYSTEM")
print("=" * 60)

# 1. Get the raw probabilities of being STABLE (Class 1)
probabilities = model.predict_proba(X_test_scaled)[: , 1]

# 2. Test multiple thresholds to find optimal trade-off
```



```
thresholds_to_test = [0.50, 0.60, 0.70, 0.75, 0.80, 0.90]

print("\nThreshold Analysis - Testing Different Safety Levels:\n")
print(f"{'Threshold':<12} {'Sensitivity':<12} {'Specificity':<12} {'Precision':<12} {'F1-Score':<12}")
print("-" * 60)

threshold_results = []

for threshold in thresholds_to_test:
    custom_pred = (probabilities >= threshold).astype(int)

    # Calculate metrics
    tn, fp, fn, tp = confusion_matrix(y_test, custom_pred).ravel()

    sensitivity = tp / (tp + fn) if (tp + fn) > 0 else 0 # True Positive Rate
    specificity = tn / (tn + fp) if (tn + fp) > 0 else 0 # True Negative Rate
    precision = tp / (tp + fp) if (tp + fp) > 0 else 0
    f1_th = 2 * (precision * sensitivity) / (precision + sensitivity) if (precision + sensitivity) > 0 else 0

    threshold_results.append({
        'Threshold': threshold,
        'Sensitivity': sensitivity,
        'Specificity': specificity,
        'Precision': precision,
        'F1': f1_th,
        'Custom_Pred': custom_pred
    })

    print(f"{'threshold':<12.2f} {'sensitivity':<12.3f} {'specificity':<12.3f} {'precision':<12.3f} {'f1_th':<12.3f}")

# 3. Recommend optimal threshold
optimal_idx = np.argmax([r['F1'] for r in threshold_results])
optimal_threshold = threshold_results[optimal_idx]['Threshold']
print(f"\n✓ RECOMMENDED THRESHOLD: {optimal_threshold:.2f} (F1-Score: {threshold_results[optimal_idx]['F1']:.4f})")

# ===== ISSUE #4: Use configurable NERC threshold from CONFIG =====
nerc_threshold = CONFIG['nerc_threshold']
custom_predictions = (probabilities >= nerc_threshold).astype(int)

print(f"\n" + "=" * 60)
```



```
print(f"NERC EARLY WARNING SYSTEM (Threshold = {nerc_threshold*100:.0f}%)の時間)")
print("=" * 60)
print("Interpretation: Grid must have ≥75% confidence to be declared SAFE")
print("Otherwise: EMERGENCY ACTION TRIGGERED\n")

nerc_cm = confusion_matrix(y_test, custom_predictions)
tn, fp, fn, tp = nerc_cm.ravel()

print("Classification Report:")
print(classification_report(y_test, custom_predictions,
                           target_names=['Emergency Action', 'Grid Safe']))

print(f"\nOperational Breakdown:")
print(f" True Negatives (Correct Emergency): {tn:5d} (avoiding false alarms)")
print(f" False Positives (False Alarms):      {fp:5d} (unnecessary activations)")
print(f" False Negatives (MISSED RISKS):      {fn:5d} ⚠ CRITICAL")
print(f" True Positives (Correct Safe):       {tp:5d}")

# 5. Visualize threshold impact
fig, axes = plt.subplots(1, 2, figsize=(15, 5))

# Plot 1: Metrics vs Threshold
thresholds_plot = [r['Threshold'] for r in threshold_results]
sensitivity_plot = [r['Sensitivity'] for r in threshold_results]
specificity_plot = [r['Specificity'] for r in threshold_results]
f1_plot = [r['F1'] for r in threshold_results]

axes[0].plot(thresholds_plot, sensitivity_plot, marker='o', label='Sensitivity (TP Rate)', linewidth=2)
axes[0].plot(thresholds_plot, specificity_plot, marker='s', label='Specificity (TN Rate)', linewidth=2)
axes[0].plot(thresholds_plot, f1_plot, marker='^', label='F1-Score', linewidth=2)
axes[0].axvline(x=nerc_threshold, color='red', linestyle='--', linewidth=2, label=f'NERC Threshold ({nerc_threshold})')
axes[0].set_xlabel('Probability Threshold')
axes[0].set_ylabel('Score')
axes[0].set_title('Threshold Optimization Analysis')
axes[0].legend()
axes[0].grid(True, alpha=0.3)

# Plot 2: NERC Confusion Matrix
```



```
sns.heatmap(nerc_cm, annot=True, fmt='d', cmap='RdYlGn',
            xticklabels=['Emergency', 'Safe'],
            yticklabels=['Actually Unstable', 'Actually Stable'],
            ax=axes[1], cbar_kws={'label': 'Count'})
axes[1].set_title(f'NERC Threshold System (Threshold={nerc_threshold})')
axes[1].set_ylabel('True Label')
axes[1].set_xlabel('System Decision')

plt.tight_layout()
plt.show()

# ===== ISSUE #10: Add prediction metadata with confidence intervals
=====
print("\n" + "=" * 60)
print("PREDICTION METADATA & CONFIDENCE ANALYSIS")
print("=" * 60)

# Create prediction dataframe with confidence info
predictions_metadata = pd.DataFrame({
    'Actual_Stability': y_test.values,
    'Predicted_Stability': y_pred,
    'Prediction_Confidence': np.maximum(y_prob, 1 - y_prob), # Confidence of the
prediction
    'Probability_Stable': y_prob,
    'Correct_Prediction': (y_test.values == y_pred).astype(int)
})

print("\nSample Predictions with Confidence Metrics:")
print(predictions_metadata.head(10).to_string())

# Analyze confidence distribution
low_conf = (predictions_metadata['Prediction_Confidence'] < 0.6).sum()
med_conf = ((predictions_metadata['Prediction_Confidence'] >= 0.6) &
            (predictions_metadata['Prediction_Confidence'] < 0.8)).sum()
high_conf = (predictions_metadata['Prediction_Confidence'] >= 0.8).sum()

print(f"\nConfidence Distribution:")
print(f"  Low Confidence (<60%): {low_conf:5d} predictions")
print(f"  Medium Confidence:      {med_conf:5d} predictions")
print(f"  High Confidence (≥80%): {high_conf:5d} predictions")

# Save metadata
metadata_path = os.path.join(CONFIG['output_dir'], 'prediction_metadata.csv')
predictions_metadata.to_csv(metadata_path, index=False)
print(f"\n✓ Prediction metadata saved to: {metadata_path}")
```

Department of electrical/electronics



```
# ===== ISSUE #2 & #5: Use CONFIG for model paths and validate before
saving =====
print("\n" + "=" * 60)
print("MODEL PERSISTENCE - SAVING FOR PRODUCTION")
print("=" * 60)

model_path = os.path.join(CONFIG['model_dir'], 'voltage_stability_model.pkl')
scaler_path = os.path.join(CONFIG['model_dir'], 'voltage_scaler.pkl')
config_path = os.path.join(CONFIG['model_dir'], 'config.pkl')

# ===== ISSUE #5: Validate models before saving =====
print("\nValidating models before persistence...")
try:
    # Test predictions
    test_pred = model.predict(X_test_scaled)
    test_prob = model.predict_proba(X_test_scaled)

    # Validate scaler
    test_scaled = scaler.transform(X_test)

    if test_pred.shape[0] > 0 and test_scaled.shape[0] > 0:
        print("✓ Model validation passed")
    else:
        raise ValueError("Model validation failed: empty predictions")
except Exception as e:
    print(f"x Model validation failed: {e}")
    raise

# Save models
joblib.dump(model, model_path)
joblib.dump(scaler, scaler_path)
joblib.dump(CONFIG, config_path)

print(f"✓ Model saved to: {model_path}")
print(f"✓ Scaler saved to: {scaler_path}")
print(f"✓ Config saved to: {config_path}")

# Save all metrics summary
metrics_summary = {
    'Model': 'Logistic Regression',
    'Test_Accuracy': accuracy,
    'Test_F1_Score': f1,
    'Test_Balanced_Accuracy': balanced_acc,
    'CV_Accuracy_Mean': cv_scores.mean(),
```

Department of electrical/electronics



```
'CV_Accuracy_Std': cv_scores.std(),
'ROC_AUC': roc_auc,
'Training_Samples': X_train.shape[0],
'Test_Samples': X_test.shape[0]
}

metrics_path = os.path.join(CONFIG['output_dir'], 'model_metrics.pkl')
joblib.dump(metrics_summary, metrics_path)
print(f"✓ Metrics saved to: {metrics_path}")

print("\nTo use the saved model in production:")
print(f" loaded_model = joblib.load('{model_path}')"")
print(f" loaded_scaler = joblib.load('{scaler_path}')"")
print(f" loaded_config = joblib.load('{config_path}')"")
print(f" X_new_scaled = loaded_scaler.transform(X_new)"")
print(f" predictions = loaded_model.predict(X_new_scaled)"")
print(f" probabilities = loaded_model.predict_proba(X_new_scaled)"")

print("\n" + "=" * 60)
print("✓ ANALYSIS COMPLETE")
print("=" * 60)
print(f"\nOutputs saved to: {CONFIG['output_dir']}/")
print(f"Models saved to: {CONFIG['model_dir']}/")
```

2. APPENDIX 2: Raw Data sample (the first 20 lines of the code)

<https://archive.ics.uci.edu/dataset/471/electrical+grid+stability+simulated+data>

```
"tau1","tau2","tau3","tau4","p1","p2","p3","p4","g1","g2","g3","g4","stab","stabf"
2.95906002455997,3.07988520422811,8.38102539191882,9.78075443222607,3.763084772063
16,-0.782603630987543,-1.25739482958732,-
1.7230863114883,0.650456460887227,0.859578105752345,0.887444920638513,0.9580339876
02737,0.0553474891727752,"unstable"
9.3040972346785,4.90252411201167,3.04754072762177,1.36935735529605,5.0678121042784
5,-1.94005842705193,-1.87274168559721,-
1.25501199162931,0.41344056837935,0.862414076352903,0.562139050527675,0.7817599106
53126,-0.00595746432603695,"stable"
8.97170690932022,8.84842842134833,3.04647874898866,1.21451813833956,3.405158180010
95,-1.20745559234302,-1.27721014673295,-
0.92049244093498,0.163041039311334,0.76688656526962,0.839444015400588,0.109853244
952427,0.00347087904838871,"unstable"
0.716414776295121,7.66959964406565,4.48664083058949,2.34056298396795,3.96379106326
633,-1.02747330413905,-1.9389441526466,-
0.997373606480681,0.446208906537321,0.976744082924302,0.929380522872661,0.36271777
426931,0.028870543444887,"unstable"
```


```

3.13411155161342,7.60877161603408,4.94375930178099,9.85757326996638,3.525810816520
96,-1.12553095451115,-1.84597485447561,-
0.554305007534195,0.797109525792467,0.455449947148291,0.656946658473716,0.82092348
6481631,0.0498603734837059,"unstable"
6.99920870915215,9.10924671643764,3.7840659093988,4.26778760341917,4.4296688228062
5,-1.85713917685067,-0.670396827477973,-
1.90213281847761,0.261792854117035,0.0779296721686504,0.542883858522323,0.46993104
1283024,-0.0173849069339299,"stable"
6.71016636762265,3.76520418352323,6.92931443867021,8.81856192345691,2.397419229379
42,-0.614589561907959,-1.20882589139371,-
0.574003776077751,0.177889561089242,0.397976580119618,0.402045952122677,0.37663027
9179848,0.0059535705684696,"unstable"
6.95351155782518,1.37912500228657,5.71940003020014,7.87030742603241,3.224495157903
25,-0.748997751247627,-1.1865169283662,-
1.28898047828942,0.371384852592056,0.633203555433068,0.732740894287254,0.380544478
575059,0.0166342974538163,"unstable"
4.68985192749584,4.00774746209312,1.47857287590383,3.73378749230253,4.041299613681
68,-1.41034431234687,-1.23820427574315,-
1.39275102559166,0.269708270661235,0.250364217202032,0.164941118523114,0.482439411
001001,-0.0386765836142913,"stable"
9.84149648830214,1.41382159178999,9.76985554870019,7.64161626330592,4.727595042018
98,-1.99136336778711,-0.857637206946418,-
1.87859446728545,0.376355980566834,0.544415315810004,0.792039424640656,0.116262822
352385,0.0123834141734175,"unstable"
5.93011011760085,6.73087337437817,6.24513790053774,0.53328751282593,2.327091592887
17,-0.702501235021046,-1.11691962966126,-
0.507670728204865,0.239815632938209,0.563109593653352,0.164461028193514,0.75370135
5880644,-0.0284114612679547,"stable"
5.38129925652915,8.01452101703768,8.0951741572488,6.76924769920496,5.5075510050451
3,-1.97271385420545,-1.84933257978301,-
1.68550457105667,0.359974025112588,0.173568816600156,0.349143850920014,0.628860423
786381,0.0281301999792468,"unstable"
1.61678683711417,2.93922810703424,0.819791186154762,4.19180384910867,3.75228191426
77,-1.48488454526459,-1.28058145097331,-
0.986815918029798,0.899698249387734,0.866545796305834,0.303920577803025,0.07761012
61113305,-0.0486167008032998,"stable"
8.55159820963844,8.31495222344878,2.54996382964303,9.92680728110744,4.891713735437
19,-1.80862626682759,-1.16706123460778,-
1.91602623400182,0.612403529208615,0.280982857520184,0.354342487890125,0.472192077
743488,0.0277556924227216,"unstable"
1.13210759360397,2.92032402330491,8.95107854465566,7.24858347796639,5.033680553347
48,-1.84607954217829,-1.36277740864708,-
1.82482360252211,0.352291990740888,0.524172882249252,0.599003855980165,0.674389868
319682,0.0148804322017331,"unstable"
7.02136199991397,4.37429363886033,4.77590415834652,8.83842594831276,3.335857461316
55,-0.962391781186988,-1.40762695274189,-
0.96583872738767,0.711099508262505,0.625363823866651,0.468334969702021,0.895143061
005827,0.0725083555585227,"unstable"
4.95224090481821,8.08867235232354,8.88331865624976,5.6945572196255,5.0672957762516
3,-1.68141419491175,-1.87706128759659,-


```

```
1.50882029374329,0.305662026508872,0.307903841937671,0.889893892795722,0.879427907
898293,0.065617204379854,"unstable"
4.1428304011016,2.43908920766944,1.29045594015738,9.45644318889404,3.9347959177254
8,-1.46929869766647,-1.766941258873,-
0.698555961186008,0.800756745514446,0.840807129209329,0.917833435936101,0.79398181
1714133,-0.00696582814861133,"stable"
9.34612594367674,7.92002984757187,2.33527630767662,3.26918132842386,4.581173917136
39,-1.10675268019072,-1.74707763942746,-
1.72734359751821,0.836076237689777,0.713253578090513,0.161517853668334,0.515983284
965043,0.0413744064662021,"unstable"
```


3. APPENDIX C: Confusion Matrix for the logistic regression model




4. APPENDIX D: ROC and AUC curves for the voltage stability machine learning model.




5. **APPENDIX E :** Feature importance for logistic regression, feature importance by absolute impact (higher means that they are more critical to decision)



6. **RECEIVER OPERATING CHARACTERISTICS CURVE:**




7. NERC THRESHOLD SYSTEM (Threshold = 0.75)




8. THRESHOLD OPTIMIZATION ANALYSIS:

This is the graph that contains sensitivity (TP Rate), Specificity (TN Rate), F1-Score, and NERC Threshold.




9. PRECISION-RECALL CURVE:

This graph shows the relationship (graphical relationship) between **Precision** (True positives) and **recall** (which is the true positive rate)



10. FEATURE CORRELATION MATRIX WITH STABILITY TARGET

Department of electrical/electronics



11. ALL OUTPUTS FOR EACH CELL OF CODE:

CELL 1

```
✓ Dataset validation passed

=====
DATA PREPROCESSING & CLASS BALANCE CHECK
=====

✓ Class Distribution BEFORE splitting:
Unstable (Class 0): 6380 samples (63.80%)
Stable (Class 1): 3620 samples (36.20%)
Training set size: 8000 samples
Testing set size: 2000 samples

=====
FEATURE CORRELATION WITH STABILITY
=====

Correlation with Stability (1=Stable, 0=Unstable):
p4      0.022785
p3      0.000649
p2     -0.006173
p1     -0.009938
g1     -0.197664
g4     -0.204931
g2     -0.217341
g3     -0.231774
tau1   -0.234898
tau3   -0.237492
tau4   -0.239375
tau2   -0.246280
dtype: float64
```

CELL 2



```
=====
DATA SCALING (for distance-based algorithms)
=====
StandardScaler centers features to mean=0, std=1
This is REQUIRED for Logistic Regression but NOT for Decision Trees

✓ Scaling complete with feature names preserved
Scaled training data shape: (8000, 12)
Scaled testing data shape: (2000, 12)
Features: ['tau1', 'tau2', 'tau3', 'tau4', 'p1', 'p2', 'p3', 'p4', 'g1', 'g2', 'g3', 'g4']
Sample scaled values (first row): [-0.424656  1.25667699  0.14023222]
```

CELL 3:

```
✓ Model trained successfully

--- CROSS-VALIDATION (5-Fold Stratified) ---
Cross-Val Accuracy: 0.8154 ± 0.0077
Cross-Val F1-Score: 0.7344 ± 0.0111
Individual fold accuracies: ['0.8100', '0.8087', '0.8269', '0.8087', '0.8225']

=====
TEST SET PERFORMANCE
=====
Accuracy:      81.85%
Balanced Accuracy: 79.25% (better for imbalanced)
F1-Score:      0.7364 (harmonic mean of precision & recall)

Detailed Classification Report:
      precision    recall  f1-score   support

Unstable (0)      0.84      0.89      0.86      1276
Stable (1)        0.78      0.70      0.74       724

   accuracy      0.82      2000
  macro avg      0.81      2000
weighted avg      0.82      2000

Confusion Matrix:
              Predicted
              Unstable  Stable
Actual Unstable    1130     146
Actual Stable       217     507
```

CELL 4

```
=====
FEATURE IMPORTANCE & GRID VULNERABILITY ANALYSIS
=====

Feature Importance Ranking:
(Positive = Stabilizing | Negative = Destabilizing)

tau4      | Coef: -0.9074 | DESTABILIZING | Impact: 0.9074
tau1      | Coef: -0.8910 | DESTABILIZING | Impact: 0.8910
tau2      | Coef: -0.8867 | DESTABILIZING | Impact: 0.8867
tau3      | Coef: -0.8657 | DESTABILIZING | Impact: 0.8657
g3        | Coef: -0.8605 | DESTABILIZING | Impact: 0.8605
g2        | Coef: -0.8129 | DESTABILIZING | Impact: 0.8129
g4        | Coef: -0.7489 | DESTABILIZING | Impact: 0.7489
g1        | Coef: -0.7301 | DESTABILIZING | Impact: 0.7301
p3        | Coef: -0.0610 | DESTABILIZING | Impact: 0.0610
p4        | Coef:  0.0405 | STABILIZING   | Impact: 0.0405
p2        | Coef:  0.0142 | STABILIZING   | Impact: 0.0142
p1        | Coef:  0.0034 | STABILIZING   | Impact: 0.0034

Model Intercept (bias): -1.0558

✓ Feature importance saved to: ./outputs/feature_importance_logistic.csv
```

CELL 5:



```
=====
DECISION TREE MODEL - COMPARISON & EVALUATION
=====
🔑 KEY DIFFERENCE: Decision Trees work on UNSCALED data (no scaling needed)

✓ Decision Tree trained on UNSCALED data

--- CROSS-VALIDATION (5-Fold Stratified) ---
Cross-Val Accuracy: 0.8455 ± 0.0116
Cross-Val F1-Score: 0.7960 ± 0.0146

=====
MODEL COMPARISON: LOGISTIC REGRESSION VS DECISION TREE
=====

Metric Logistic Regression Decision Tree
Accuracy          81.85%          85.85%
Balanced Accuracy  79.29%          85.50%
F1-Score           0.7364          0.8117
CV Accuracy Mean   0.8154          0.8455

✓ BEST MODEL: Decision Tree (85.85% accuracy)
```

CELL 6:

```
--- DECISION TREE FEATURE IMPORTANCE ---

Feature Importance
3 tau4 0.153021
2 tau3 0.146177
1 tau2 0.145699
0 tau1 0.128405
8 g1 0.115032
```

CELL 7:

```
0.60 0.612 0.922 0.816 0.699
0.70 0.507 0.956 0.868 0.640
0.75 0.463 0.969 0.896 0.610
0.80 0.388 0.979 0.912 0.545
0.90 0.233 0.997 0.977 0.377

✓ RECOMMENDED THRESHOLD: 0.50 (F1-Score: 0.7364)

=====
NERC EARLY WARNING SYSTEM (Threshold = 75%)の時間
=====
Interpretation: Grid must have ≥75% confidence to be declared SAFE
Otherwise: EMERGENCY ACTION TRIGGERED

Classification Report:
precision recall f1-score support
Emergency Action 0.76 0.97 0.85 1276
Grid Safe 0.90 0.46 0.61 724

accuracy 0.79 2000
macro avg 0.83 0.72 0.73 2000
weighted avg 0.81 0.79 0.76 2000

Operational Breakdown:
True Negatives (Correct Emergency): 1237 (avoiding false alarms)
False Positives (False Alarms): 39 (unnecessary activations)
False Negatives (MISSED RISKS): 389 ⚠️ CRITICAL
True Positives (Correct Safe): 335
```

CELL 8

Part 1

```
=====
PREDICTION METADATA & CONFIDENCE ANALYSIS
=====

Sample Predictions with Confidence Metrics:
  Actual_Stability Predicted_Stability Prediction_Confidence Probability_Stable Correct_Prediction
0                0                0          0.926760          0.073240                1
1                0                0          0.976694          0.023306                1
2                0                0          0.968499          0.031501                1
3                1                1          0.730231          0.730231                1
4                0                0          0.501402          0.498598                1
5                1                1          0.843976          0.843976                1
6                1                1          0.511465          0.511465                1
7                0                0          0.992617          0.007383                1
8                0                0          0.968502          0.031498                1
9                1                1          0.833186          0.833186                1

Confidence Distribution:
Low Confidence (<60%):    236 predictions
Medium Confidence:       549 predictions
High Confidence (≥80%):  1215 predictions

✓ Prediction metadata saved to: ./outputs/prediction\_metadata.csv

=====
MODEL PERSISTENCE - SAVING FOR PRODUCTION
=====

Validating models before persistence...
```

Part 2

```
Medium Confidence:    549 predictions
High Confidence (≥80%): 1215 predictions

✓ Prediction metadata saved to: ./outputs/prediction\_metadata.csv

=====
MODEL PERSISTENCE - SAVING FOR PRODUCTION
=====

Validating models before persistence...
✓ Model validation passed
✓ Model saved to: ./models/voltage\_stability\_model.pkl
✓ Scaler saved to: ./models/voltage\_scaler.pkl
✓ Config saved to: ./models/config.pkl
✓ Metrics saved to: ./outputs/model\_metrics.pkl

To use the saved model in production:
loaded_model = joblib.load('./models/voltage_stability_model.pkl')
loaded_scaler = joblib.load('./models/voltage_scaler.pkl')
loaded_config = joblib.load('./models/config.pkl')
X_new_scaled = loaded_scaler.transform(X_new)
predictions = loaded_model.predict(X_new_scaled)
probabilities = loaded_model.predict_proba(X_new_scaled)

=====
✓ ANALYSIS COMPLETE
=====

Outputs saved to: ./outputs/
Models saved to: ./models/
```

And that's a wrap!